

The Oracle logo is displayed in a bold, red, sans-serif typeface. The letters are thick and blocky, with a distinctive design where the 'R' and 'A' have a diagonal cut. A registered trademark symbol (®) is positioned at the top right of the 'E'.

ORACLE®

“This presentation is for informational purposes only and may not be incorporated into a contract or agreement.”

This document is for informational purposes. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described in this document remains at the sole discretion of Oracle. This document in any form, software or printed matter, contains proprietary information that is the exclusive property of Oracle. This document and information contained herein may not be disclosed, copied, reproduced or distributed to anyone outside Oracle without prior written consent of Oracle. This document is not part of your license agreement nor can it be incorporated into any contractual agreement with Oracle or its subsidiaries or affiliates.

THE INFORMATION COMPANY

Andrew Holdsworth

Director of Real World Performance

Mohamed Ziauddin

Hal Takahara

Sunil Chakkappen

Optimizer Development Team

ORACLE

A Practical Approach to Optimizer Statistics in Oracle Database 10g

Agenda

- Background to the subject of Optimizer Statistics
- Proposed methodology
- The General Solution
- Refining The Gathering Process
- Dealing with Exceptions
- Rule to Cost Based Migrations
- Dealing with Volatile Tables
- Questions & Answers

Background to the Subject of Optimizer Statistics

- The need for a methodology
- Sorry DBAs, it will mean more work
- Best practices are still evolving
- The payback from good statistics management and execution plans will exceed any benefit of init.ora tuning by orders of magnitude

Proposed Methodology 2005

- Start with the General Solution (90% of SQL)
- Refine the method based upon actual data shape and usage (next 9% of SQL)
- Correct the last queries (remaining 1% of SQL)
- Data Dictionary tables also need statistics from 10g onwards

The General Solution

- Build Statistics using the following method:

- Use dbms_stats package not analyze command
- In 9i

```
DBMS_STATS.gather_schema_stats(  
  ownname=>'<schema>',  
  estimate_percent=>dbms_stats.auto_sample_size  
  cascade=>TRUE,  
  method_opt=>'FOR ALL COLUMNS SIZE AUTO')
```

- In 10g statistics get gathered automatically

Refining/Adapting the Gathering Process

- When to gather statistics ?
- Dealing with evolving data shapes and sizes
 - Choosing Sampling Sizes and Types
 - DML monitoring at runtime
 - Column Usage History Information
 - Partitioned Objects Issues

When to Gather Statistics

- Candidates
 - After large amounts of data change (loads, purges, bulk updates)
 - New high/low values for keys generated
 - Newly created tables
 - Upgrading CPUs, I/O subsystem (system statistics)
 - RBO to CBO migrations
 - New database creations
- Apply common sense to rebuild frequency

Sample Sizes

- Default recommendation
 - `estimate_percent => DBMS_STATS.AUTO_SAMPLE_SIZE`
 - This is great for uniform sizes and distributions of data
 - Issues can occur with skewed data or if there are lot of nulls in the column
- Smaller sample size for large objects
 - Should save time
 - Row sampling (Full scan of data sampling rows)
 - Block sampling (Reduced I/O)
 - Same issues as for `AUTO_SAMPLE_SIZE`

DML Monitoring

- Used by `dbms_stats` to identify objects with "stale" statistics
- On by default in 10g, not in 9i
`alter table <table_name> monitoring;`
- Tracked in
[DBA|ALL|USER]_TAB_MODIFICATIONS
- 9i and 10g use a 10% change as the threshold to recalculate stats

DML Monitoring Example

```
SQL> select table_name, partition_name, inserts, updates, deletes
       from all_tab_modifications;
```

TABLE_NAME	PARTITION_NAME	INSERTS	UPDATES	DELETES
-----	-----	-----	-----	-----
USERS	USERS14	0	3328	0
ACCOUNT	ACCOUNT23	23450	738	230
EMP	EMP19	2409	390	0

Column Usage History

- Used by `dbms_stats` to identify candidate columns on which to build histograms
- Column Usage Dictionary Table
- Lets Talk About About Histograms

Column Usage Dictionary Table

- Tracked in `sys.col_usage$`
- Captures predicate column name and access method type
- Used by DBMS_STATS to gather histograms when `method_opt=>'for ... size auto'`
- Data is persisted after database shutdown
- This is why the stats built after running SQL **can** change/evolve

Column Usage Example

TAB_NAME	COLUMN_NAME	EQUALITY	EQUIJOIN	NONEQUIJOIN	RANGE	LIKE	NULL
ACCOUNT	ORGANIZATION_ID	26	0	0	0	0	2
ACCOUNT	ACCOUNT_ID	26	25	0	0	0	1
ACCOUNT	OWNER	0	25	0	0	0	1
ACCOUNT	DELETED	26	0	0	0	0	1
ACC_SHARE	ORGANIZATION_ID	27	0	0	0	0	2
ACC_SHARE	ACCOUNT_ID	0	25	0	0	0	2
ACC_SHARE	UG_ID	0	26	0	0	0	2
ACC_SHARE	DELETED	27	0	0	0	0	2

* SQL for query is in slide notes

Lets Talk About Histograms

- Generally help with skewed data
- Should be used with literals in statements
- `method_opt =>`
 - `'for all columns'` defaults to 75 buckets
 - `'for all columns size auto'` defaults to 254 buckets **
- Issues can arise when there are many unique values with skewed data

Partitioned Objects

- The option for local statistics exists to allow partition level optimization
- Issues arise when running with large numbers of partitions/sub partitions > 1000
- Dropping these objects/stats takes too long to be practical
- Examples include the composite range/hash partitioned table of 36 months x 32 hash subpartitions (1152 total partitions)
- Exchanging partitions in/out may/may not invoke stats gathering.

Dealing with the Exceptions

- Init.ora parameters
- Hints
- SQL profiles
- Tough Decisions
- Things statistics cannot deduce

Init.ora Parameters

- The most common forms of abuse
 - optimizer_index_caching
 - optimizer_index_cost_adj
 - db_file_multiblock_read_count
- The issue here is moving forward and impact on other applications

Use of Hints

- Used to fine tune queries
 - Incorporate application knowledge not possible from the collected stats
 - Dealing with impossible complex SQL
 - Dealing with volatile or dynamic tables
 - Dealing with and working around bugs
 - Dealing with column correlations
- Again the issue is moving forward to new releases

Use of SQL profiles

- Contains tuning information used by the optimizer
- Introduced in 10g – `dbms_sqltune`
- More exhaustive analysis than CBO alone
- Store SQL profile to get tuned plan for future executions

Managing Statistics

- In 9i use `DBMS_STATS.[EXPORT|IMPORT]*STATS` procedures to back up and restore statistics or to move the statistics between databases
- In 10g
 - Use `DBMS_STATS.[EXPORT|IMPORT]*STATS` procedures to move statistics between databases or to experiment with different sets of statistics
 - Use `DBMS_STATS.RESTORE*STATS` procedures to recover older versions of statistics. There are also procedures to manage the retention and purging of statistics histories.
- Controlling when Statistics get re-built
 - In 9i `DBMS_STATS` package to rebuild statistics
 - In 10g statistics gets gathered automatically.
Use `DBMS_SCHEDULER.DISABLE('GATHER_STATS_JOB')` to disable automatic statistics gathering.

RBO to CBO Migration Issues

- The importance of testing and running of workload
- Dealing with very complicated schemas
 - Complicated SQL(Joins, RBO predicates, encode application knowledge)
 - Objects with large number of access paths

Dealing with Volatile Tables

- Basic Premise
 - Good plans with large data sets are better than sub-optimal plans with small datasets.
- Building Stats vs. Dynamic Sampling (leave table without stats but lock it with)
- Seeding Statistics (store representative table stats then lock it)
- Lock stats with
`dbms_stats.lock_table_stats`

Detection when Good Statistics yield poor plans

- The most common reason for poor execution plans with perceived “good” statistics is inaccurate row count estimates
 - This leads to bad access path selection
 - This leads to bad join method selection
 - This leads to bad join order selection
- In summary one bad row count estimate can cascade into a very poor plan

Determining the Bad Source

- Basic Analysis:
 1. Is the rows estimate from explain plan correct?
 2. Is the rows estimate correct in all_[tab|part]_col_statistics?
 3. Is the sample size large enough to give accurate row estimates?
 4. Obtain accurate stats and recheck the plan

Bad Source Example

- Stats gathered with sample of `dbms_stats.auto_sample_size`
- Data is known to be skewed and have high number of distinct values(NDV)

Bad Source Example cont.

```
select a.name,
       a.account_id,
       a.site,
       a.address1_state,
       a.phone1,
       a.type_enum,
       u.alias,
       u.users_id,
       upper (a.name),
from   sales.account a,
       (select distinct s.account_id
        from             core.ug_blowout b,
                        sales.acc_share s
        where             s.organization_id = '00D3000000002jy' and
                        b.organization_id = '00D3000000002jy' and
                        b.users_id = '005300000000dQfh' and
                        s.ug_id = b.ug_id and
                        s.deleted = '0') t,
       core.users u
where  (t.account_id = a.account_id) and
       (u.users_id = a.owner) and
       (a.account_id != '0000000000000000') and
       (a.organization_id = '00D3000000002jy' and
        a.deleted = '0') and
       (u.organization_id = '00D3000000002jy')
```

Bad Source Example cont.

Original Plan from v\$sql_plan_statistics_all

Query time - Elapsed: 00:02:51.05

Statistics from v\$sql_plan_statistics_all

CARDINALITY = optimizer row estimate

LAST_OUTPUT_ROWS = actual row count

OPERATION	OBJECT_NAME	CARDINALITY	LAST_OUTPUT_ROWS
-----	-----	-----	-----
VIEW		21,249	1,087
SORT		21,249	1,087
HASH JOIN		21,249	2,172
TABLE ACCESS USERS	USERS	625	3,326
INDEX IEUSERSACTIVE	IEUSERSACTIVE	1	3,326
HASH JOIN		63,120	2,172
HASH JOIN		293,480	2,172
INDEX AKUG_BLOWOUT	AKUG_BLOWOUT	73	13
INDEX IEACC_SHARE_DIVISION	IEACC_SHARE_DIVISION	3,899,279	14,030,696
TABLE ACCESS ACCOUNT	ACCOUNT	370,639	1,609,704

* SQL for above query is in slide notes

Bad Source Example cont.

- What was wrong?
 - Row estimates are off causing wrong plan
 - Comparison of all_part_col_statistics and actual distinct values showed the NDV estimate was significantly off due to data skew
- How to fix it (choose one)?
 - Re-gather stats with compute or a high enough sample to give good estimates
 - Use dbms_stats.set_column_stats to set the column distinct correctly

Bad Source Example cont.

Bad Sample Estimate Comparison

Numbers taken from all_part_col_statistics and
“count(distinct <column_name>)” for given partition

<u>TABLE NAME</u>	<u>COLUMN NAME</u>	<u>ESTIMATED NDV</u>	<u>ACTUAL NDV</u>
ACCOUNT	ORGANIZATION_ID	325	1,907
ACCOUNT	ACCOUNT_ID	2,952,246	2,963,757
ACCOUNT	OWNER	1,857	8,963
ACC_SHARE	ORGANIZATION_ID	71	306
ACC_SHARE	UG_ID	969	8,991
ACC_SHARE	DELETED	2	2
USERS	ORGANIZATION_ID	2,527	1,906
USERS	USERS_ID	2,985	27,342
UG_BLOWOUT	ORGANIZATION_ID	2,071	1,552
UG_BLOWOUT	USERS_ID	20,252	19,783

Bad Source Example cont.

New Plan from v\$sql_plan_statistics_all

Query time - Elapsed: 00:00:00.61

Statistics from v\$sql_plan_statistics_all

CARDINALITY = optimizer row estimate

LAST_OUTPUT_ROWS = actual row count

OPERATION	OBJECT_NAME	CARDINALITY	LAST_OUTPUT_ROWS
-----	-----	-----	-----
VIEW		1009	1087
SORT		1009	1087
NESTED LOOPS		2019	2172
NESTED LOOPS		2019	2172
NESTED LOOPS		2019	2172
INDEX AKUG_BLOWOUT	AKUG_BLOWOUT	6	13
TABLE ACCESS ACC_SHARE	ACC_SHARE	2019	2172
INDEX AK2ACC_SHARE	AK2ACC_SHARE	2041	2190
TABLE ACCESS ACCOUNT	ACCOUNT	2019	2172
INDEX PKACCOUNT	PKACCOUNT	2019	2172
TABLE ACCESS USERS	USERS	2019	2172
INDEX PKUSERS	PKUSERS	2019	2172

A large, stylized graphic of the letters 'Q' and 'A' in a dark grey, serif font. A large, vibrant red ampersand is positioned between the 'Q' and 'A', partially overlapping them. The text 'QUESTIONS' and 'ANSWERS' is written in white, bold, sans-serif capital letters across the middle of the graphic.

QUESTIONS ANSWERS

